

Matthew Steffen

Leader with talents for communication and consensus-building. Interested in both Manager and I.C. roles.

📍 San Carlos, CA ✉ m@tthew.io ☎ 713-208-6899 in m12n 🌐 msteffen

Skills

- **Engineering:** Extensive professional experience with C++, Java, and Go codebases, plus smaller projects in Typescript+React and Python. I am always excited to learn new languages and tools.
- **Management:** Managed up to five engineers and co-managed seven, and I have hired into backend and full-stack roles. I have strong written and verbal communication skills and wrote many of Pachyderm's internal templates, handbooks and guides.

Experience

Hewlett Packard Enterprise (via Pachyderm acquisition)

*San Jose & Remote
March 2023 – Nov 2024*

Software Engineering Manager

- I was designated a Key Employee when HPE acquired Pachyderm in 2023.
- Continued managing the Pachyderm Integrations Team through May 2024, shipping a complete rewrite of our Jupyter extension to meet the security requirements of HPE's customers.
- Co-managed a team of seven engineers, after HPE re-organized Pachyderm into this single team in June 2024. To align the disparate teams being combined, I rewrote our engineering handbook, overhauled team processes, and worked with HPE IT to implement a common development environment. Overall efficiency improved, and we shipped several long-backlogged features.

Pachyderm

*San Francisco & Remote
Jan 2022 – Feb 2023*

Software Engineering Manager

- Hired for and supported the team through the design and development of our JupyterLab extension.
- This extension, our biggest project, was a developer tool that allowed data scientists to explore and consume huge Pachyderm datasets directly in Jupyter's UI (which required deep modifications to Jupyter's frontend) and run Jupyter notebooks as sharded, orchestrated Pachyderm pipelines without modification.
- Contributed to design reviews and stayed in our on-call rotation throughout my tenure as manager.

Software Engineer


Aug 2016 – Dec 2021

- Pachyderm is a distributed filesystem with versioning, for data and AI pipelines. When Pachyderm runs a pipeline, its pipeline orchestrator uses its versioning metadata to processes only new and updated shards. Results from unchanged shards are re-used, cutting the runtime and cost of pipelines by 90% or more, in some cases. Sometimes described as "Bazel for data".
- Implemented several major features, including an authorization system and a complete rewrite of our pipeline controller, among others (Go, Kubernetes).
- Started new Integrations Team, to connect Pachyderm and other AI/ML products (Kubeflow, Seldon, LabelStudio, AzureML, Spark, and others). I acted as tech lead through its growth to five engineers.

Google

*Mountain View, CA
Aug 2011 – July 2016*

Software Engineer

- Independently maintained a legacy authorization service used by Google Docs, Photos, YouTube, Search, and other products through 2012 (Java).
- Joined the [Zanzibar](#)  team in 2013 as the fourth engineer. I migrated users from the legacy service to Zanzibar and implemented features and performance improvements in Zanzibar itself (C++).
- Contributed features and performance improvements to Cloud IAM in late 2015 and 2016.

Education

University of Chicago

Sept 2007 – May 2011

BS in Computer Science (also met requirements for BS in Mathematics) - 3.8 GPA

- ICPC World Finalist (2010, 2011) and USA Powerlifting Collegiate National Finalist (2011)